

Generating human faces with Adversarial Networks

Hai Bacti

January 14, 2021

© research.nvidia.com

This time we'll train a neural net to generate plausible human faces in all their subtlety: appearance, expression, accessories, etc. 'Cuz when us machines gonna take over Earth, there won't be any more faces left. We want to preserve this data for future iterations. Yikes...

Based on <https://github.com/Lasagne/Recipes/pull/94>.

```
[1]: import sys
      sys.path.append("..")
      import grading
      import download_utils
      import tqdm_utils
```

```
[2]: import matplotlib.pyplot as plt
      %matplotlib inline
      import numpy as np
      plt.rcParams.update({'axes.titlesize': 'small'})

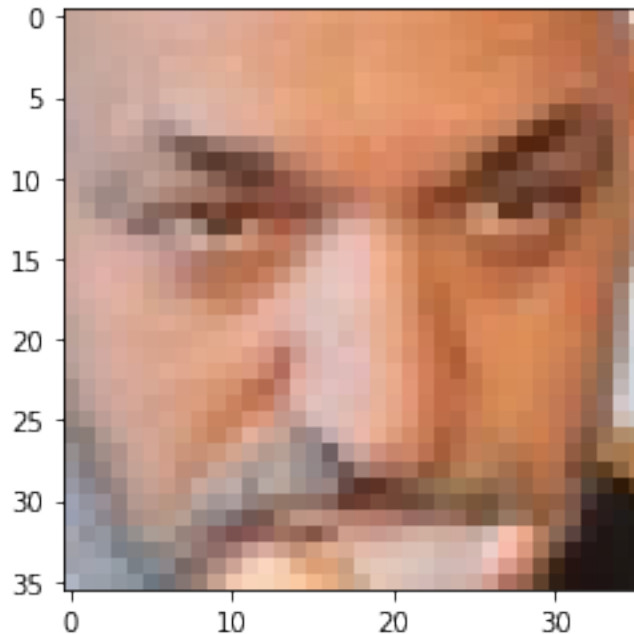
      from sklearn.datasets import load_digits
      #The following line fetches you two datasets: images, usable for autoencoder
      ↪ training and attributes.
      #Those attributes will be required for the final part of the assignment
      ↪ (applying smiles),
      #so please keep them in mind
      from lfw_dataset import load_lfw_dataset
      data, attrs = load_lfw_dataset(dimx = 36, dimy = 36)

      #preprocess faces
      data = np.float32(data) / 255.
      IMG_SHAPE = data.shape[1:]
```

```
0%|          | 0/13233 [00:00<?, ?it/s]
```

```
[3]: #print random image
      plt.imshow(data[np.random.randint(data.shape[0])], cmap = "gray", interpolation
      ↪ = "none")
```

```
[3]: <matplotlib.image.AxesImage at 0x7fd1f98d7d50>
```



1 Generative adversarial nets 101

© *torch.github.io*

Deep learning is simple, isn't it? * build some network that generates the face (small image) * make up a **measure of how good that face is** * optimize with gradient descent :)

The only problem is: how can we engineers tell well-generated faces from bad? And i bet you we won't ask a designer for help.

If we can't tell good faces from bad, we delegate it to yet another neural network!

That makes the two of them: * **G**enerator - takes random noise for inspiration and tries to generate a face sample. * Let's call him **G(z)**, where z is a gaussian noise. * **D**iscriminator - takes a face sample and tries to tell if it's great or fake. * Predicts the probability of input image being a **real face** * Let's call him **D(x)**, x being an image. * **D(x)** is a prediction for real image and **D(G(z))** is prediction for the face made by generator.

Before we dive into training them, let's construct the two networks.

```
[4]: import tensorflow as tf
      from tensorflow.keras.models import Sequential
      from tensorflow.keras import layers as L
```

```
[5]: CODE_SIZE = 256

      generator = Sequential()
```

```

generator.add(L.InputLayer([CODE_SIZE], name = 'noise'))
generator.add(L.Dense(10 * 8 * 8, activation = 'elu'))

generator.add(L.Reshape((8, 8, 10)))
generator.add(L.Conv2DTranspose(64, kernel_size = (5, 5), activation = 'elu'))
generator.add(L.Conv2DTranspose(64, kernel_size = (5, 5), activation = 'elu'))
generator.add(L.UpSampling2D(size = (2, 2)))
generator.add(L.Conv2DTranspose(32, kernel_size = 3, activation = 'elu'))
generator.add(L.Conv2DTranspose(32, kernel_size = 3, activation = 'elu'))
generator.add(L.Conv2DTranspose(32, kernel_size = 3, activation = 'elu'))

generator.add(L.Conv2D(3, kernel_size = 3, activation = None))

```

```

[6]: assert generator.output_shape[1:] == IMG_SHAPE, \
      "generator must output an image of shape %s, but instead it produces %s" % (
      IMG_SHAPE, generator.output_shape[1:])

```

1.0.1 Discriminator

- Discriminator is your usual convolutional network with interlooping convolution and pooling layers
- The network does not include dropout/batchnorm to avoid learning complications.
- We also regularize the pre-output layer to prevent discriminator from being too certain.

```

[7]: discriminator = Sequential()
discriminator.add(L.InputLayer(IMG_SHAPE))
discriminator.add(L.Conv2D(filters = 32, kernel_size = (3, 3), padding = 'same', activation = 'elu'))
discriminator.add(L.MaxPooling2D(pool_size = (2, 2)))
discriminator.add(L.Conv2D(filters = 64, kernel_size = (3, 3), padding = 'same', activation = 'elu'))
discriminator.add(L.MaxPooling2D(pool_size = (2, 2)))
discriminator.add(L.Conv2D(filters = 128, kernel_size = (3, 3), padding = 'same', activation = 'elu'))
discriminator.add(L.MaxPooling2D(pool_size = (2, 2)))
discriminator.add(L.Conv2D(filters = 256, kernel_size = (3, 3), padding = 'same', activation = 'elu'))
discriminator.add(L.MaxPooling2D(pool_size = (2, 2)))
discriminator.add(L.Flatten())
discriminator.add(L.Dense(256, activation = 'tanh'))
discriminator.add(L.Dense(2, activation = tf.nn.log_softmax))

```

2 Training

We train the two networks concurrently: * Train **discriminator** to better distinguish real data from **current** generator * Train **generator** to make discriminator think generator is real * Since discriminator is a differentiable neural network, we train both with gradient descent.

© *deeplearning4j.org*

Training is done iteratively until discriminator is no longer able to find the difference (or until you run out of patience).

2.0.1 Tricks:

- Regularize discriminator output weights to prevent explosion
- Train generator with **adam** to speed up training. Discriminator trains with SGD to avoid problems with momentum.
- More: <https://github.com/soumith/ganhacks>

```
[8]: def discriminator_training(real_data, noise):
    disc_optimizer.minimize \
        (loss = lambda: d_loss(real_data, noise), var_list = discriminator.
         ↪ trainable_weights)

    def d_loss(real_data, noise):
        logp_real = discriminator(real_data)
        logp_gen = discriminator(generator(noise))
        loss = -tf.reduce_mean(logp_real[:, 1] + logp_gen[:, 0])
        loss += tf.reduce_mean(discriminator.layers[-1].kernel ** 2) # regularize
        return loss

    disc_optimizer = tf.keras.optimizers.SGD(learning_rate = 1e-3)
```

```
[9]: def generator_training(noise):
    gen_optimizer.minimize(loss = lambda: g_loss(noise), var_list = generator.
    ↪ trainable_weights)

    def g_loss(noise):
        logp_gen = discriminator(generator(noise))
        return -tf.reduce_mean(logp_gen[:, 1])

    gen_optimizer = tf.keras.optimizers.Adam(learning_rate = 1e-4)
```

2.0.2 Auxiliary functions

Here we define a few helper functions that draw current data distributions and sample training batches.

```
[10]: def sample_noise_batch(bsize):
        return np.random.normal(size = (bsize, CODE_SIZE)).astype('float32')

def sample_data_batch(bsize):
    idxs = np.random.choice(np.arange(data.shape[0]), size = bsize)
    return data[idxs]

def sample_images(nrow,ncol, sharp = False):
    images = generator.predict(sample_noise_batch(bsize = nrow * ncol))
    if np.var(images) != 0:
        images = images.clip(np.min(data), np.max(data))
    for i in range(nrow * ncol):
        plt.subplot(nrow, ncol, i + 1)
        if sharp:
            plt.imshow(images[i].reshape(IMG_SHAPE), cmap = "gray",
↪ interpolation = "none")
        else:
            plt.imshow(images[i].reshape(IMG_SHAPE), cmap = "gray")
    plt.show()

def sample_probas(bsize):
    plt.title('Generated vs real data')
    plt.hist (
        np.exp(discriminator.predict(sample_data_batch(bsize)))[:, 1],
        label = 'D(x)', alpha = 0.5, range = [0, 1]
    )
    plt.hist (
        np.exp(discriminator.predict(generator.
↪ predict(sample_noise_batch(bsize))))[:, 1],
        label = 'D(G(z))', alpha = 0.5, range = [0, 1]
    )
    plt.legend(loc = 'best')
    plt.show()
```

2.0.3 Training

Main loop. We just train generator and discriminator in a loop and plot results once every N iterations.

```
[13]: from IPython import display

num_epochs = 10000
num_batch = 100

for epoch in tqdm_utils.tqdm_notebook_failsafe(range(num_epochs)):
    real_data = sample_data_batch(num_batch)
```

```

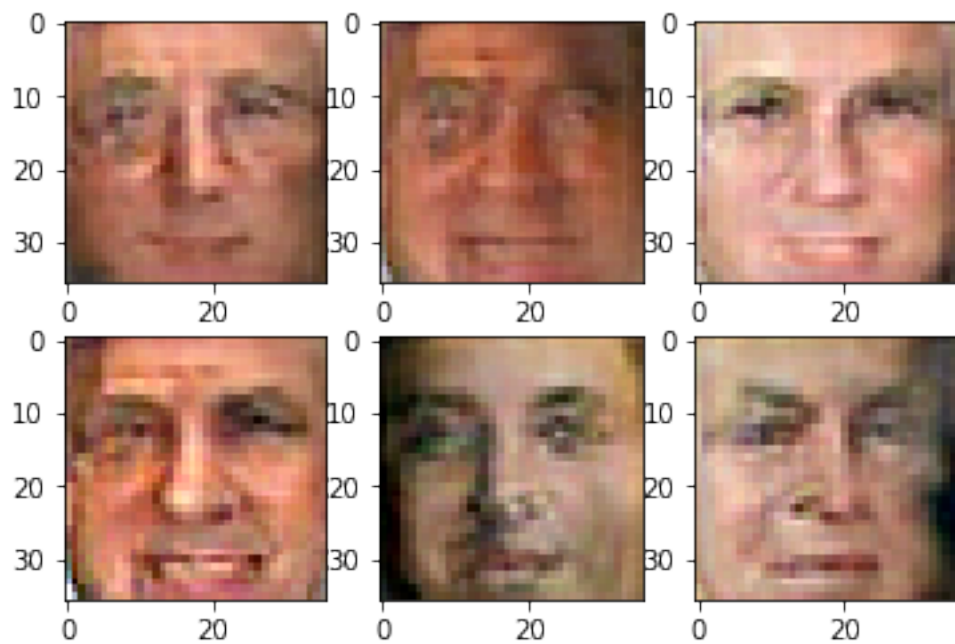
noise = sample_noise_batch(num_batch)
print("Epoch {:05d}".format(epoch))


for i in range(5):
    discriminator_training(real_data, noise)

generator_training(noise)

if epoch % 100 == 0:
    display.clear_output(wait = True)
    sample_images(2, 3, True)
    sample_probas(1000)

```





```
[14]: #The network was trained for about 15k iterations.  
#Training for longer yields MUCH better results  
plt.figure(figsize = [16, 24])  
sample_images(16, 8)
```

